

# Lecture 18

## Restoring Invariants

15-122: Principles of Imperative Computation (Spring 2018)  
Frank Pfenning

In this lecture we will implement heaps and operations on them. The theme of this lecture is reasoning with invariants that are partially violated, and making sure they are restored before the completion of an operation. We will only briefly review the algorithms for inserting and deleting the minimal node of the heap; you should read the notes for the previous lecture on priority queues and keep them close at hand.

**Computational Thinking:** We convert ideas developed diagrammatically in the last lecture into working code.

**Algorithms and Data Structures:** Temporarily violating and restoring invariants is a common theme in algorithms. It is a technique you need to master.

**Programming:** We practice writing generic code involving arrays.

### 1 The Heap Structure

We use the following header struct to represent heaps.

```
typedef struct heap_header heap;
struct heap_header {
    int limit;                // limit = capacity+1
    int next;                // 1 <= next && next <= limit
    elem[] data;             // \length(data) == limit
    has_higher_priority_fn* prior; // != NULL
};
```

The field `prior` is provided by the client and tells us how to compare elements.

```
// f(x,y) returns true if e1 has STRICTLY higher priority than e2
typedef bool has_higher_priority_fn(elem e1, elem e2);
```

Since the significant array elements start at 1, as explained in the previous lecture, the `limit` must be one greater than the desired capacity. To prevent overflows, we will also require that `limit` be less than `int_max()/2`. The next index must be between 1 and `limit`, and the element array must have exactly `limit` elements.

## 2 Minimal Heap Invariants

Before we implement the operations, we define a function that checks the heap invariants. The shape invariant is automatically satisfied due to the representation of heaps as arrays, but we need to carefully check the ordering invariants. It is crucial that no instance of the data structure that is not a true heap will leak across the interface to the client, because the client may then incorrectly call operations that require heaps with data structures that are not.

First, we check that the heap is not `NULL`, that `limit` has acceptable values, and that the length of the array matches the given `limit`. The latter must be checked in an annotation, because, in C0 and C1, the length of an array is not available to us at runtime except in contracts. Second, we check that `next` is in range, between 1 and `limit`. Finally, we check that the client-provided comparison is defined and non-`NULL`.

```
1 bool is_heap_safe(heap* H) {
2   return H != NULL
3     && (1 < H->limit && H->limit < int_max()/2)
4     && is_array_expected_length(H->data, H->limit)
5     && (1 <= H->next && H->next <= H->limit)
6     && H->prior != NULL;
7 }
```

This is not sufficient to know that we have a valid heap! The specification function `is_heap_safe` is the minimal specification function we need to be able to access the data structure; we want to make sure anything we pass to the user additionally satisfies the ordering invariant.

This invariant acts as the precondition of some of our helper functions. We first use the client's function, accessible as `H->prior`, to express a more

useful concept for our implementation: that the element in index  $i$  can be correctly placed as the parent of the element in index  $j$  in the heap.

```

9  bool ok_above(heap* H, int i, int j)
10 //@requires is_heap_safe(H);
11 //@requires 1 <= i && i < H->next;
12 //@requires 1 <= j && j < H->next;
13 {
14     return !(*H->prior)(H->data[j], H->data[i]);
15 }

```

This function returns true whenever the element at position  $i$  has priority higher than or equal to the element at position  $j$ .

A second helper function that uses `is_heap_safe` swaps an element with its parent:

```

17 void swap_up(heap* H, int i)
18 //@requires is_heap_safe(H);
19 //@requires 2 <= i && i < H->next;
20 {
21     elem tmp = H->data[i];
22     H->data[i] = H->data[i/2];
23     H->data[i/2] = tmp;
24 }

```

### 3 The Heap Ordering Invariant

It turns out to be simpler to specify the ordering invariant in the second form seen in the last lecture, which stipulates that each node except the root needs to be greater or equal to its parent. To check this we iterate through the array and compare the priority of each node `data[i]` with its parent, except for the root ( $i = 1$ ) which has no parent.

```

26 bool is_heap_ordered(heap* H)
27 //@requires is_heap_safe(H);
28 {
29     for (int i = 2; i < H->next; i++) {
30         //@loop_invariant 2 <= i;
31         if (!ok_above(H, i/2, i)) return false;
32     }
33     return true;
34 }
35

```

```
36 bool is_heap(heap* H) {  
37     return is_heap_safe(H) && is_heap_ordered(H);  
38 }
```

Observe that the loop starts at index 2, since the root of the heap is stored at index 1.

## 4 Creating Heaps

We start with the simple code to test if a priority queue implemented as a heap is empty or full, and to allocate a new (empty) heap. A heap is empty if the next element to be inserted would be at index 1. A heap is full if the next element to be inserted would be at index `limit` (the size of the array).

```
1 bool pq_empty(heap* H)  
2 //@requires is_heap(H);  
3 {  
4     return H->next == 1;  
5 }  
6  
7 bool pq_full(heap* H)  
8 //@requires is_heap(H);  
9 {  
10    return H->next == H->limit;  
11 }
```

To create a new heap, we allocate a struct and an array and set all the right initial values.

```
13 heap* pq_new(int capacity, has_higher_priority_fn* prior)  
14 //@requires 0 < capacity && capacity < int_max()/2 - 1;  
15 //@requires prior != NULL;  
16 //@ensures is_heap(\result) && pq_empty(\result);  
17 {  
18     heap* H = alloc(heap);  
19     H->limit = capacity+1;  
20     H->next = 1;  
21     H->data = alloc_array(elem, H->limit);  
22     H->prior = prior;  
23     return H;  
24 }
```

Note that `H->data[0]` is unused. We could have allocated an array of exactly capacity at the cost of complicating our index operations. The precondition `capacity < int_max()/2 - 1` wards against overflows.

## 5 Insert and Sifting Up

The shape invariant tells us exactly where to insert the new element: at the index `H->next` in the data array. Then we increment the next index.

```

25 void pq_add(heap* H, elem x)
26 //@requires is_heap(H) && !pq_full(H);
27 //@ensures is_heap(H);
28 {
29   H->data[H->next] = x;
30   (H->next)++;           // basic invariants hold
31   // but ordering invariant may be violated
32   // ...
33 }
```

By inserting  $x$  in its specified place, we have, of course, violated the ordering invariant. We need to *sift up* the new element until we have restored the invariant. The invariant is restored when the new element is bigger than or equal to its parent or when we have reached the root. We still need to sift up when the new element is less than its parent. This suggests the following code:

```

32   int i = H->next - 1;           // element we just added
33   while (i > 1 && !ok_above(H,i/2,i)) {
34     swap_up(H, i);
35     i = i/2;
36   }
```

Setting `i = i/2` is moving up in the tree, to the place we just swapped the new element to.

At this point, as always, we should ask why accesses to the elements of the priority queue are safe. By short-circuiting of conjunction, we know that `i > 1` when we ask whether `H->data[i/2]` is okay above `H->data[i]`. But we need a loop invariant to make sure that it respects the upper bound. The index `i` starts at `H->next - 1`, so it should always be strictly less than `H->next`.

```

33   while (i > 1 && !ok_above(H,i/2,i))
```

```

34  //@loop_invariant 1 <= i && i < H->next;
35  {
36      swap_up(H, i);
37      i = i/2;
38  }

```

One small point regarding the loop invariant: we just incremented  $H \rightarrow next$ , so it must be strictly greater than 1 and therefore the invariant  $1 \leq i$  must be satisfied.

But how do we know that swapping the element up the tree restores the ordering invariant? We need an additional loop invariant which states that  $H$  is a valid heap *except at index  $i$* . Index  $i$  may be smaller than its parent, but it still needs to be less than or equal to its children. We therefore postulate a function `is_heap_except_up` and use it as a loop invariant.

```

33  while (i > 1 && !ok_above(H,i/2,i))
34  //@loop_invariant 1 <= i && i < H->next;
35  //@loop_invariant is_heap_except_up(H, i);

```

The next step is to write this function. We copy the `is_heap_ordered` function, but check a node against its parent only when it is different from the distinguished element where the exception is allowed. The differences are highlighted.

```


bool is_heap_except_up(heap* H, int n)
//@requires is_heap_safe(H);
//@requires 1 <= n && n < H->next;
{
    for (int i = 2; i < H->next; i++) {
        //@loop_invariant 2 <= i;
        if (!(i == n || ok_above(H, i/2, i))) return false;
    }
    return true;
}

```

We observe that `is_heap_except_up(H, 1)` is equivalent to `is_heap(H)`. That's because the loop over  $i$  starts at 2, so the exception  $i \neq n$  is always true.

Now we try to prove that this is indeed a loop invariant, and therefore our function is correct. Rather than using a lot of text, we verify this properties on general diagrams. Other versions of this diagram are entirely symmetric. On the left is the relevant part of the heap before the swap and

on the right is the relevant part of the heap after the swap. The relevant nodes in the tree are labeled with their priority. Nodes that may be above  $a$  or below  $c, c_1, c_2$  and to the right of  $a$  are not shown. These do not enter into the invariant discussion, since their relations between each other and the shown nodes remain fixed. Also, if  $x$  is in the last row the constraints regarding  $c_1$  and  $c_2$  are vacuous.




We know the following properties on the left from which the properties shown on the right follow as shown:

<i>Current iteration</i>			<i>Next iteration</i>		
$a \leq b$	(1)	order	$a ? x$	allowed exception	
$b \leq c$	(2)	order	$x \leq c$	from (5) and (2)	
$x \leq c_1$	(3)	order	$x \leq b$	from (5)	
$x \leq c_2$	(4)	order	$b \leq c_1$	??	
$b > x$	(5)	since we swap	$b \leq c_2$	??	

(For this and similar examples, we'll assume that we're using a min-heap.) Our invariant gives us no way to know that  $b \leq c_1$  and  $b \leq c_2$ . We see that simply stipulating the (temporary) invariant that every node is greater or equal to its parent except for the one labeled  $x$  is not strong enough. It is not necessarily preserved by a swap.

But we can strengthen it a bit. You might want to think about how before you move on to the next page.

The strengthened invariant also requires that the children of the potentially violating node  $x$  are greater or equal to their grandparent! Let's reconsider the diagrams.



We have more assumptions on the left now ((6) and (7)), but we have also two additional proof obligations on the right ( $a \leq c$  and  $a \leq b$ ).

Current iteration			Next iteration		
$a \leq b$	(1)	order	$a ? x$	allowed exception	
$b \leq c$	(2)	order	$a \leq c$	from (1) and (2)	
$x \leq c_1$	(3)	order	$a \leq b$	(1)	
$x \leq c_2$	(4)	order	$x \leq c$	from (5) and (2)	
$b > x$	(5)	since we swap	$x \leq b$	from (5)	
$b \leq c_1$	(6)	grandparent	$b \leq c_1$	(6)	
$b \leq c_2$	(7)	grandparent	$b \leq c_2$	(7)	

Success! We just need an additional function that checks this loop invariant:

```
bool grandparent_check(heap* H, int n)
//@requires is_heap_safe(H);
//@requires 1 <= n && n < H->next;
{
    if (n == 1) return true;           // Reached the root
    if (n*2 >= H->next) return true;   // No children
    if (n*2 + 1 == H->next)            // Left child only
        return ok_above(H, n/2, n*2);
    return ok_above(H, n/2, n*2)      // Both children
        && ok_above(H, n/2, n*2 + 1);
}
```



Using this additional invariant, we have a loop that provably restores the `is_heap` invariant.

```

32 while (i > 1 && !ok_above(H,i/2,i))
33   //@loop_invariant 1 <= i && i < H->next;
34   //@loop_invariant is_heap_except_up (H, i);
35   //@loop_invariant grandparent_check(H, i);
36   {
37     swap_up(H, i);
38     i = i/2;
39   }

```

Note that the strengthened loop invariants (or, rather, the strengthened definition of what it means to be a heap except in one place) is not necessary to show that the postcondition of `pq_add` (i.e., `is_heap(H)`) is implied.

**Postcondition:** If the loop exits, we know the loop invariants and the negated loop guard are true:

$1 \leq i < next$  (LI 1)

`is_heap_except_up(H, i)` (LI 2)

Either  $i \leq 1$  or `ok_above(H, i/2, i)` Negated loop guard

We distinguish the two cases.

**Case:**  $i \leq 1$ . Then  $i = 1$  from (LI 1), and `is_heap_except_up(H, 1)`. As observed before, that is equivalent to `is_heap(H)`.

**Case:** `ok_above(H, i/2, i)`. Then the only possible index  $i$  where `is_heap_except_up(H, i)` makes an exception and does not check whether `ok_above(H, i/2, i)` is actually no exception, and we have `is_heap(H)`.

Overall, the function `pq_add` is as follows

```

26 void pq_add(heap* H, elem e)
27 //@requires is_heap(H) && !pq_full(H);
28 //@ensures is_heap(H);
29 {
30   H->data[H->next] = e;
31   (H->next)++;
32   /* H may no longer be a heap! */
33
34   int i = H->next - 1;
35   while(i > 1 && !ok_above(H,i/2,i))
36     //@loop_invariant 1 <= i && i < H->next;
37     //@loop_invariant is_heap_except_up(H, i);
38     //@loop_invariant grandparent_check(H, i);
39     {
40       swap_up(H, i);
41       i = i/2;
42     }
43 }
```

## 6 Deleting the Minimum and Sifting Down

Recall that deleting the minimum swaps the root with the last element in the current heap and then applies the *sifting down* operation to restore the invariant. As with insert, the operation itself is rather straightforward, although there are a few subtleties. First, we have to check that  $H$  is a heap, and that it is not empty. Then we save the minimal element, swap it with the last element (at `next - 1`), and delete the last element (now the element that was previously at the root) from the heap by decrementing `next`.

```

76 elem pq_rem(heap* H)
77 //@requires is_heap(H) && !pq_empty(H);
78 //@ensures is_heap(H);
79 {
80   elem min = H->data[1];
81   (H->next)--;
82
83   if (H->next > 1) {
84     H->data[1] = H->data[H->next]; // Swap last element in
```


```

85     // Ordering invariant may be violated
86     sift_down(H);
87 }
88 return min;
89 }

```

Next we need to restore the heap invariant by sifting down from the root, with `sift_down(H)`. We only do this if there is at least one element left in the heap.

But what is the precondition for the sifting down operation? Again, we cannot express this using the functions we have already written. Instead, we need a function `is_heap_except_down(H, n)` which verifies that the heap invariant is satisfied in  $H$ , except possibly at  $n$ . This time, though, it is between  $n$  and its children where things may go wrong, rather than between  $n$  and its parent as in `is_heap_except_up(H, n)`. In the pictures below this would be at  $n = 1$  on the left and  $n = 2$  on the right.



We change the test accordingly.

```

/* Valid heap except at n, looking down the tree */
bool is_heap_except_down(heap H, int n)
//@requires is_heap_safe(H);
//@requires 1 <= n && n < H->next;
{
    for (int i = 2; i < H->next; i++) {
        //@loop_invariant 2 <= i;
        if (!(i/2 == n || ok_above(H, i/2, i))) return false;
    }
    return true;
}

```

```
void sift_down(heap H)
//@requires is_heap_safe(H);
//@requires H->next > 1 && is_heap_except_down(H, 1);
//@ensures is_heap(H);
{
    int i = 1;

    while (2*i < H->next)
        //@loop_invariant 1 <= i && i < H->next;
        //@loop_invariant is_heap_except_down(H, i);
        //@loop_invariant grandparent_check(H, i);
        // ...
}
```

We want to return from the function if we have restored the invariant, that is if the element in index  $i$  is okay above all of his children. However, there may be either 1 or 2 children (the loop guard checks that there will be at least one). So we have to guard this access by a bounds check. Clearly, when there is no right child, checking the left one is sufficient.

```
while (2*i < H->next)
//@loop_invariant 1 <= i && i < H->next;
//@loop_invariant is_heap_except_down(H, i);
//@loop_invariant grandparent_check(H, i);
{
    int left = 2*i;
    int right = left+1;

    if (ok_above(H, i, left)           // All good on the left, and
        && (right >= H->next ||         // no right child or
            ok_above(H, i, right)))    // all good on the right too
        return;                       // Nothing to do
    // ...
}
```

If this test fails, we have to determine the smaller of the two children. If there is no right child, we pick the left one, of course. Once we have found the smaller one we swap the current one with the smaller one, and then make the child the new current node  $i$ .

Here is the overall function:

```
45 void sift_down(heap H)
46 //@requires is_heap_safe(H);
47 //@requires H->next > 1 && is_heap_except_down(H, 1);
48 //@ensures is_heap(H);
49 {
50     int i = 1;
51
52     while (2*i < H->next)
53         //@loop_invariant 1 <= i && i < H->next;
54         //@loop_invariant is_heap_except_down(H, i);
55         //@loop_invariant grandparent_check(H, i);
56     {
57         int left = 2*i;
58         int right = left+1;
59
60         if (ok_above(H, i, left)           // All good on the left, and
61             && (right >= H->next ||        // no right child or
62                 ok_above(H, i, right))) // all good on the right too
63             return;                       // Nothing to do!
64         if (right >= H->next ||            // No right child, or
65             ok_above(H, left, right)) { // left is smaller or equal
66             swap_up(H, left);
67             i = left;
68         } else {                          // right is smaller
69             //@assert right < H->next && ok_above(H, right, left);
70             swap_up(H, right);
71             i = right;
72         }
73     }
74
75     //@assert i < H->next && 2*i >= H->next;
76     //@assert is_heap_except_down(H, i);
77     return;
78 }
```

Before the second return, we know that `is_heap_except_down(H, i)` and  $2i \geq \text{next}$ . This means there is no node  $j$  in the heap such that  $j/2 = i$  and the exception in `is_heap_except_down` will never apply.  $H$  is indeed a heap.

At this point we should give a proof that `is_heap_except_down` is really an invariant. This is left as Exercise 4.

## 7 Heapsort

We rarely discuss testing in these notes, but it is useful to consider how to write decent test cases. Mostly, we have been doing random testing, which has some drawbacks but is often a tolerable first cut at giving the code a workout. It is *much* more effective in languages that are type safe such as C0, and even more effective when we dynamically check invariants along the way.

In the example of heaps, one nice way to test the implementation is to insert a random sequence of numbers, then repeatedly remove the minimal element until the heap is empty. If we store the elements in an array in the order we take them out of the heap, the array should be sorted when the heap is empty! This is the idea behind heapsort. We first show the code, using the random number generator we have used for several lectures now, then analyze the complexity. As the priority function, we use `int_lt(x, y)` which returns `true` if and only if  $x < y$ . The standard loop invariants have been omitted for conciseness.

```
int main() {
    int n = (1<<9)-1;           // 1<<9 for -d; 1<<13 for timing
    int num_tests = 10;         // 10 for -d; 100 for timing
    int seed = 0xc0c0ffee;
    rand_t gen = init_rand(seed);
    int[] A = alloc_array(int, n);
    heap H = pq_new(n, &int_lt);

    print("Testing heap of size "); printint(n);
    print(" "); printint(num_tests); print(" times\n");
    for (int j = 0; j < num_tests; j++) {
        for (int i = 0; i < n; i++) {
            pq_add(H, rand(gen));
        }
        for (int i = 0; i < n; i++) {
```

```
    A[i] = pq_rem(H);
}
assert(pq_empty(H));           // heap not empty
assert(is_sorted(A, 0, n));    // heapsort failed
}
print("Passed all tests!\n");
return 0;
}
```

Now for the complexity analysis. Inserting  $n$  elements into the heap is bounded by  $O(n \log n)$ , since each of the  $n$  inserts is bounded by  $\log n$ . Then the  $n$  element deletions are also bounded by  $O(n \log n)$ , since each of the  $n$  deletions is bounded by  $\log n$ . So altogether we get  $O(2n \log n) = O(n \log n)$ . Heapsort is asymptotically as good as mergesort or as good as the expected complexity of quicksort with random pivots.

The sketched algorithm uses  $O(n)$  auxiliary space, namely the heap. One can use the same basic idea to do heapsort in place, using the unused portion of the heap array to accumulate the sorted array.

Testing, including random testing, has many problems. In our context, one of them is that it does not test the strength of the invariants. For example, say we write no invariants whatsoever (the weakest possible form), then compiling with or without dynamic checking will always yield the same test results. We really should be testing the invariants themselves by giving examples where they are not satisfied. However, we should not be able to construct such instances of the data structure on the client side of the interface. Furthermore, within the language we have no way to “capture” an exception such as a failed assertion and continue computation.

## 8 Summary

We briefly summarize key points of how to deal with invariants that must be temporarily violated and then restored.

1. Make sure you have a clear high-level understanding of why invariants must be temporarily violated, and how they are restored.
2. Ensure that at the interface to the abstract type, only instances of the data structure that satisfy the full invariants are being passed. Otherwise, you should rethink all the invariants.

3. Write predicates that test whether the partial invariants hold for a data structure. Usually, these will occur in the preconditions and loop invariants for the functions that restore the invariants. This will force you to be completely precise about the intermediate states of the data structure, which should help you a lot in writing correct code for restoring the full invariants.

## Exercises

**Exercise 1.** Write a recursive version of `is_heap_ordered`.

**Exercise 2.** Write a recursive version of `is_heap_except_up`.

**Exercise 3.** Write a recursive version of `is_heap_except_down`.

**Exercise 4.** Give a diagrammatic proof for the invariant property of sifting down for delete (called `is_heap_except_down`), along the lines of the one we gave for sifting up for insert.

**Exercise 5.** Above, we separated out the sift down operation into its own function `sift_down`. Do the same for sift up.

**Exercise 6.** Say we want to extend priority queues so that when inserting a new element and the queue is full, we silently delete the element with the lowest priority (= maximal key value) before adding the new element. Describe an algorithm, analyze its asymptotic complexity, and provide its implementation.

**Exercise 7.** Using the invariants described in this lecture, write a function `heapsort` which sorts a given array in place by first constructing a heap, element by element, within the same array and then deconstructing the heap, element by element.

[Hint: It may be easier to sort the array in descending order and reverse in a last pass or use so called max heaps where the maximal element is at the top.]

**Exercise 8.** Is the array `H->data` of a heap always sorted?